

# Quick to production with the best of Spark and Tensorflow

Ronny Mathew

Manager, Data Science

*Rue Gilt Groupe*





# Who are we

# RUE GILT G R O U P E

**Rue Gilt Groupe** is a fashion technology company based in Boston

Our **50M+ members** get exclusive access to **private sales** on **must have brands** from an ever changing catalog of **3M+ products**

[ruelala.com](http://ruelala.com)



[gilt.com](http://gilt.com)

# Data Science at Rue Gilt Groupe (RGG)


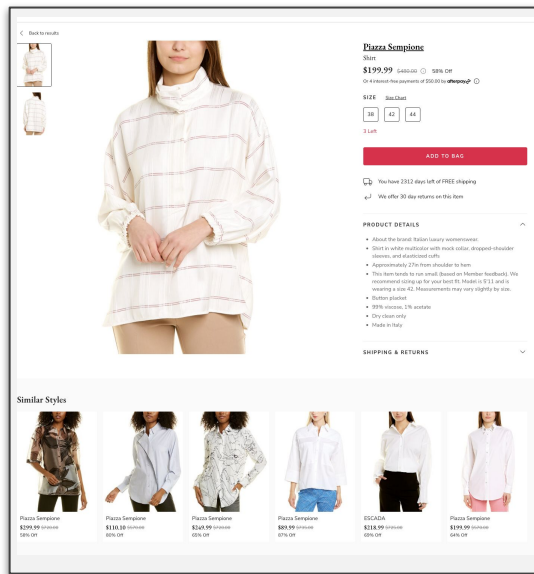
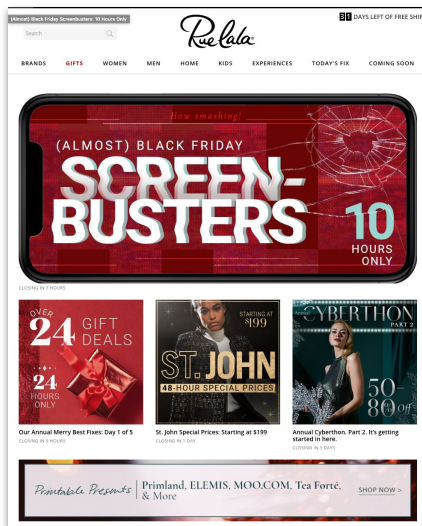
RUE GILT  
GROUPE

Close-knit group of passionate data scientists who thrive on collaboration and innovation

Working on various DS projects including NLP, CV, time-series Forecasting, Recommenders

We use ML to **personalize** the site and email experiences of our members among others projects

Sorting  
main  
pages



Product recs

# Our 2020 ML stack

We leverage Spark and container based ML pipelines extensively

In addition to this, we use common ML packages like sklearn, xgBoost, openCV, prophet, ...



# Developing next-gen Models

*Developing the next generation of products pushed us to expand **beyond Spark ML** and leverage state-of-the-art deep learning models in **Tensorflow***

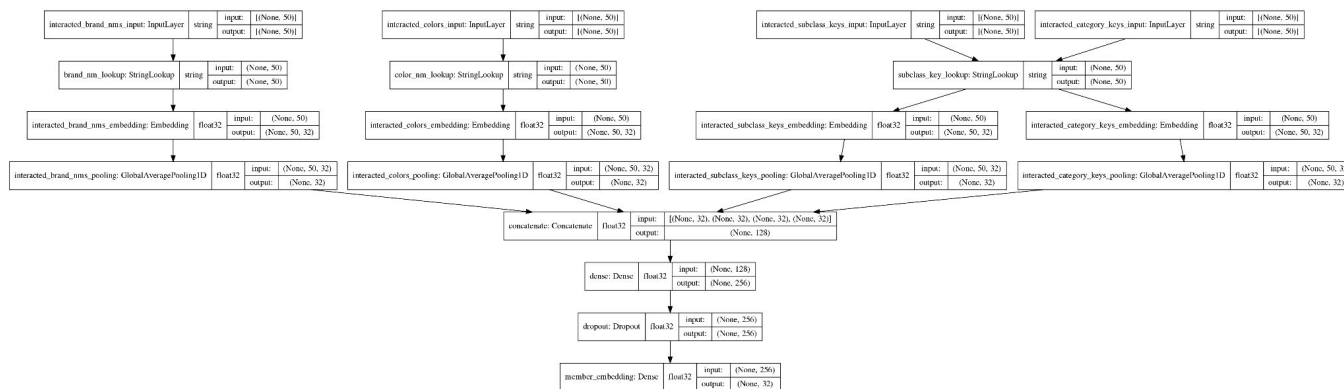
- Solution to **quickly** move from a working prototype to a production model
- Did not want added complexity with more MLOps code
- **Minimal** approach perfect for a **small-medium** full stack Data science team
- Enable **rapid prototyping and experimentations**
- Lets Data Scientists and ML Engineers **focus more on research** and **actual ML tasks**

# What we will cover

- This talk will focus on how we were able to achieve all of the above using some recent advancements in Tensorflow 2 and extension libraries for Spark by merging the **best of both Spark and Tensorflow worlds**
- **Enable Quick Distributed training** in tensorflow on GPU and CPU clusters
- MLFlow model registry as a **central model storage** ensuring consistency for all downstream batch and realtime inference jobs

# How this helped us at RGG

- Models using huge clickstream datasets are difficult to train on a single instance
- Developed a prototype for a **GRU based personalization model for new members**
- Converted our prototype to a production model quickly in **less than 2 weeks**
- A/B Test proved our new members loved their deep learning recs with **>25% lift in CTR!**



# Our 2021 next-gen ML stack

Using the tools we will cover today, we successfully integrated Tensorflow into our spark world

Flexibility to choose from Spark or Tensorflow ecosystem at any ML stage





# Next up for us

- Develop deep learning models for all of our personalization and recommenders
- Online session based recommendations
- Better text representation with cutting edge language models (BERT and more)
- Product Tagging and Catalog Management with NER, text and image models
- Leverage image resources better with Autoencoders and vision models
- Online recommendations using Reinforcement Learning
- Realtime intent detection for a user session



**Let's get started with  
some Spark features**



# Apache Spark

- Popular open source big data processing platform
- Write code in multiple programming languages
- Run pipelines on a single-node or clusters
- Work on structured and unstructured data
- Batch and streaming pipelines
- Distributed Machine learning with Spark ML



# Data Pipelines in multiple languages



- Easy-to-use APIs in SQL, Python, Java, R, and/or Scala
- User-defined functions extend spark SQL functionalities

## Data engineers

```
# in python
import pyspark.sql.functions as f
@f.udf("string")
def say_hello(name: str) -> str:
    return f"Hello {name}"
sqlContext.udf.register("say_hello",
say_hello)
```

## Data Analysts/Data Scientists

```
-- uses the udf in sql
SELECT say_hello('Bob')
```

# Pandas UDFs



Allows using a pandas UDF to manipulate spark dataframes

Manipulate spark dataframe as pandas DataFrames instead of a per row transformation

```
@f.pandas_udf("string")
def simple_udf(iterator: Iterator[pd.Series]) -> Iterator[pd.Series]:
    for x in iterator:
        yield pd.Series(list(map(lambda r: r + "1", x)))

spark.range(10).select(simple_udf("id")).show()
```

[Read more about them here](#)

- Scalable implementations of many common ML algorithms, including
  - Traditional ML models like Regression, Trees, Multilayer Perceptron, Naive Bayes
  - Text models like Bag of words, Word2Vec, TF-IDF, LDA
  - ALS-based Collaborative Filtering and FP-Growth mining
  - Locality Sensitive Hashing and Random Projection for nearest neighbor lookups
- Libraries are available to easily integrate sklearn, XGBoost, tensorflow, pytorch and others



**Let's look at some  
Tensorflow features**



# Tensorflow

- Machine learning framework maintained and used within Google
- V2 has adopted the core principle of **progressive disclosure of complexity** from Keras

*You should always be able to get into lower-level workflows in a gradual way. You shouldn't fall off a cliff if the high-level functionality doesn't exactly match your use case. You should be able to gain more control over the small details while retaining a commensurate amount of high-level convenience.*

[source](#)

- TFX extends tensorflow beyond model building and to end to end ML pipelines
- A lot of extension libraries that extends Tensorflow capabilities
- Great tutorials and guides on the official website for anyone to get started quickly
  - <https://www.tensorflow.org/tutorials>
  - <https://www.tensorflow.org/guide>





- TF APIs for data loading and manipulation using Tensorflow datasets
- Recommended way to load data into a tensorflow 2.x model
- **TfRecordDataset** is a tf dataset that can load data directly from **tfRecord** files
- Does a lot of great optimizations behind the scenes including prefetching and caching
  - While training on GPU, prefetch uses the CPU to load the next n batches
  - **tf.data.AUTOTUNE** can automatically adapt prefetching (and others) to your cluster

```
dataset = (tf.data.TfRecordDataset(filename=[list of tfrecord files]))  
          .batch(1000)  
          .shuffle(10000)  
          .prefetch(tf.data.AUTOTUNE)
```

# TF Distribute



- Provides distributed training strategies in synchronous/asynchronous mode, on GPU/TPU, etc
- Automatically assigns **default strategy** to every training loop
- A single-node development code can be converted to a distributed code by adding a few lines initially when we start development

```
# we can also add more conditions here, eg check for num workers
# add more strategies as needed
if tf.config.list_physical_devices('GPU'):
    my_strategy = tf.distribute.MirroredStrategy()
else: # Use the Default Strategy
    my_strategy = tf.distribute.get_strategy()

with my_strategy:
    # wrap all the training code within this scope
    ...
```

# Distributing Datasets



Convert a tf dataset to a distributed dataset in a few lines of code

```
# There are two shard policies, DATA and FILE
options = tf.data.Options()
options.experimental_distribute.auto_shard_policy =
tf.data.experimental.AutoShardPolicy.DATA
dataset.with_options(options)
```

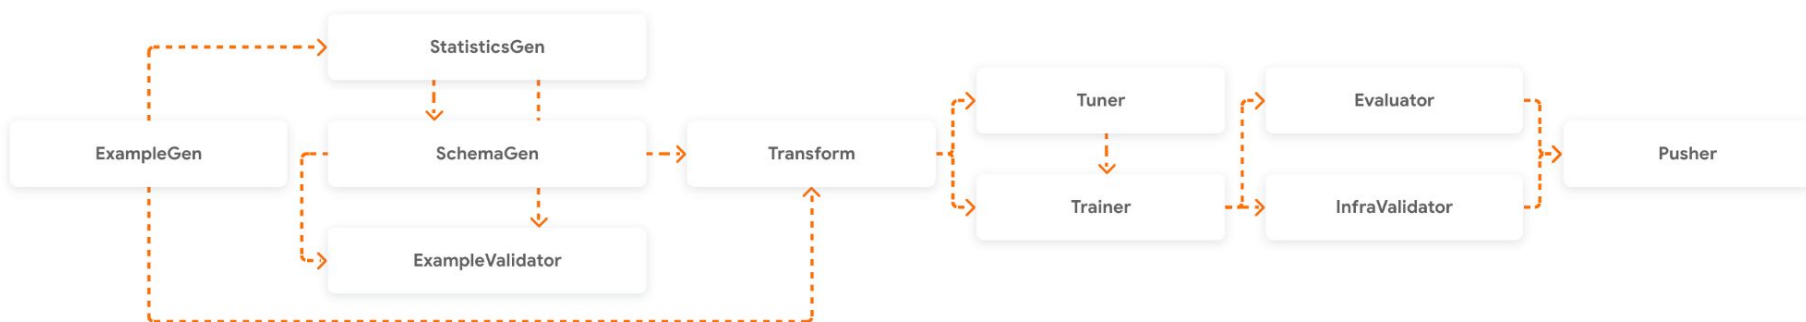
# TF Extended (TFx)

- End-to-end machine learning platform for tensorflow
- Everything from data ingestion to model serving and monitoring

Ingest & validate data

Train & analyze model

Deploy in production



ML Metadata →



TensorFlow Data  
Validation →



TensorFlow Transform  
→



TensorFlow →



TensorFlow Model  
Analysis →



TensorFlow Serving, TF  
Lite & TFJS →

# Some cool TFX components



- **Tensorflow Serving**
  - Core Tensorflow component responsible for deploying and online serving
  - Loads and serves a model at a high performant REST endpoint
  - Docker containers provided by google with different flavors of tensorflow
  - Can be scaled as needed using dockerized containers for online recommendations
- **Tensorflow Transform:**
  - Performs feature engineering and transforms
  - Alternative to using Spark (as in our case)
- **Tensorflow Model Analysis (TFMA)**
  - Model validations
  - Monitoring issues and drifts in models

# TF Recommenders



- Extension library for building Deep learning Recommenders
- Provides easy to use wrappers for models, custom metrics and losses
- Exposes google's *scann* approximate nearest neighbor engine as a Layer
- Simplifies building models for
  - Retrieval (narrow down entire corpus to a few thousand candidates)
  - Ranking (rank top candidates to a handful with more complex architecture)
  - Multitask recommenders (learn from multiple objectives)
  - Deep Cross Networks (capture explicit feature interactions between items)
  - and more

# More Extension Libraries and tools



- **TF Ranking**
  - Build Learning-to-Rank (LTR) models
  - commonly used ranking loss functions like pointwise, pairwise, and listwise
  - ranking metrics like Mean Reciprocal Rank and NDCG
- **TF Agents**
  - Enables building contextual bandit and RL agents
- **TF Hub**
  - Collection of the latest pre-trained models for text, image, video, audio, and more
  - Easily pull to any existing model as a layer with simple code,

```
# sometimes we also need to load preprocessing layers
bert_path="https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/4"
bert_embeddings = hub.KerasLayer(bert_path, trainable=False)
```



# MLFlow



- Open source platform for ML lifecycle
- Enables Experiment tracking, reproducibility, deployability and model management
- Works with most common ML libraries
- Seamlessly integrates with Spark and tensorflow

## MLflow Tracking

Record and query experiments: code, data, config, and results

[Read more](#)

## MLflow Projects

Package data science code in a format to reproduce runs on any platform

[Read more](#)

## MLflow Models

Deploy machine learning models in diverse serving environments

[Read more](#)

## Model Registry

Store, annotate, discover, and manage models in a central repository

[Read more](#)

# MLFlow Tracking



- Track Hyperparameters, Metrics and Artifacts associated with each and every run
- **Compare metrics** from different runs and use that to **select** the best model
- **Autologging** capabilities for common frameworks

## ▼ Parameters

Name	Value
cross val horizon	14 days
cross val initial	730 days
cross val period	10 days
cross validation end	03/21/2021
cross validation start	03/22/2017
prophet changepoint_prior_scale	0.05
prophet mcmc_samples	0
prophet seasonality_mode	multiplicative
train months	25

## ▼ Metrics

Name	Value
<a href="#">coverage</a>	0.581
<a href="#">cross val avg coverage</a>	0.733
<a href="#">cross val avg mae</a>	74229.3
<a href="#">cross val avg mape</a>	0.086
<a href="#">cross val avg mse</a>	9557433561.4
<a href="#">cross val avg rmse</a>	97349.4
<a href="#">horizon 10 days coverage</a>	0.82
<a href="#">horizon 10 days mae</a>	67163.5
<a href="#">horizon 10 days mape</a>	0.072

## ▼ Artifacts

- ▶  config
- ▶  data
- ▶  model\_runs

# MLFlow Model Registry



- **Manage models** using MLFlow models and Model Registry
- Transition model versions between Staging, Production and Archived stage
- Single source of truth for all models and its versions ensuring **consistency**

## Registered Models

Name 	Latest Version	<div>Staging</div>	<div>Production</div>
Forecast Test	Version 1	—	—
pricing_demand_model	Version 84	Version 33	Version 84
pricing_demand_model_backtest	Version 70	Version 70	Version 16
residual_pricing_demand_model	Version 182	—	Version 174
residual_pricing_demand_model_backtest	Version 182	Version 182	—



# Linking them all

# TF records with Spark

- TFRecord is a tf native file format that stores sequences of binary records as **Examples**
- Reading and writing TFRecords from Spark ensures that Spark and TFX uses same format
- Enables the flexibility of using Spark or TFX at any point of our ML lifecycle
- Other approaches like Petastorm achieves similar goals but loses the interoperability with TFX

```
# Write tfRecords as Example from a Spark Dataframe
some_df.write.format("tfrecords")
    .option("recordType", "Example")
    .mode("overwrite")
    .save(save_path)
```

# Spark Tensorflow Distributor


- Part of the TensorFlow ecosystem
- Enables running **tf.distribute strategies** from inside Spark jobs
- **Delegates** cluster management and communication to **Spark**
- Also runs on CPU clusters
- Straightforward usage,
  - Define both dataset loading and TF model building code in a **train** method
  - Pass the number of workers in the cluster and train method to the distributed runner

```
num_workers = 10
# add local_mode=True for local testing
# you can also specify another strategy in the train method and turn on
custom_strategy=True
MirroredStrategyRunner(num_slots=num_workers).run(train)
```

# Spark Tensorflow Distributor

Under the hood this library automates a lot of infrastructure code

- Communication between nodes
- designation of chief and worker nodes
- Supports custom\_strategy to use other tf.distribute strategies (TPUs, GPUs, etc)
  - Update train method with new strategy, add custom\_strategy=True to the runner



# Who is the chief here

- At times we need to know what the role of a node is, especially which one is chief
- Treat callbacks and other operations differently on chief
  - Don't need all the workers (over)writing the model checkpoints to the same path!
- `tf.distribute` works by specifying a `tf_config` environment variable on each node
  - config also includes a node index where *node 0 is the chief node*
- We can find the role of a node using this code,

```
if 'TF_CONFIG' in os.environ:  
    tf_config = json.loads(os.environ['TF_CONFIG'])  
    node_index = tf_config['task']['index']  
    is_chief = node_index == 0
```



# Dont forget to save your models!

Since the training happens on the workers and not just a single instance,

**we don't have access to the final trained model unless we handle it explicitly!**

That means you might be running training for days *but there would be no way to retrieve your final model (yikes!)*.

We can use ModelCheckpoint or BackupAndRestore callbacks to handle this

```
save_best_callback = ModelCheckpoint(  
    filepath=s3://persist_path, # Somewhere on persistent  
    storage  
    monitor=monitor metric,  
    save_best_only=True)
```

# Batch Inference in Spark

- Pull the latest model from model registry
- Load it in a pandas UDF
- Define a schema for the UDF output
- Simply use the new UDF in your inference pipeline as you would use any spark function
- We can also use MLFlow function to create a UDF for a model with `mlflow.pyfunc.spark_udf`

# Recap of our new pipeline

<i>ML Stage</i>	<i>How</i>
<b>Data Ingestion</b>	Use Spark to read from data sources
<b>Feature Engineering</b>	Create data pipelines in Spark Write data to persistent store as TFRecords
<b>Model Building</b>	Read TFRecords in distributed code from persistent store Train distributed tensorflow model on Spark Save trained model to MLFlow
<b>Model Selection and Management</b>	Pick the best model from MLFlow and add to model registry Use the same production model in all downstream applications
<b>Batch Inference</b>	Load model from MLFlow registry into a pandas UDF Add the UDF into the inference Spark pipeline

# Unlocking Realtime Inference

- Realtime inference enables exciting use cases involving direct feedback from our members
- Leverage Spark Streaming pipelines to process the realtime data
- Build an online inference service using TF serving and the model from model registry
- Some use cases are,
  - In-session Recommendations
  - Intent and shopping pattern detection
  - Location based trending products
  - Hyper personalization using Reinforcement Learning

# Pros of this Approach

- Great for smaller teams who want to focus on ML and less infrastructure
- Quickly transition from prototype to production models
- Flexibility to choose right framework for the use case at any stage of the ML lifecycle
- Scalable deep learning pipelines with minimal MLOps code
- Run distributed deep learning loads on Spark GPU/CPU clusters with fault tolerance
- Use all of tensorflow ecosystem and extension libraries
- Central storage, versioning, and management of models using MLflow

# Open Challenges

- Distributed training doesn't handle fault recovery automatically (in the case of a dead worker)
  - Parameter server strategy could be a better fit for this case
- Tuning and finding the right amount of nodes for the cluster still requires some manual effort
- Cryptic and messy Tensorflow error messages
- APIs between Keras and Tensorflow are in a state of flux

# Q&A

Thank you!